

Fault Localization



Objective

- Explore the existing tools and methods for fault localization.
- Identify the exact location of the contended locks.

Perf tool

We can see the symbols from the memory spaces which are consuming the highest percentage of the CPU. This happens because multiple threads are running creating lock contention.

```
1 0.01% psi_task_change
1 0.01% amd_pmu_disable_all
1 0.01% newidle_balance
1 0.01% schedule
1 0.01% perf_event_task_tick
1 0.01% timekeeping_advance
1 0.01% update_load_avg
1 0.01% interrupt_entry
1 0.01% load_balance
1 0.01% x86_pmu_disable_all
1 0.01% update_dl_rq_load_avg
1 0.01% find_next_and_bit
1 0.01% __sched_text_start
1 0.01% irq_work_run_list
1 0.01% task_tick_fair
1 0.01% sched_clock_cpu
2 0.02% __update_load_avg_se
2 0.02% sleep_micros
2 0.02% update_nohz_stats
2 0.02% _raw_spin_lock_irqsave
2 0.02% MM_CompactSchemeFixupObject::verifyForwardingPtr
2 0.02% __update_load_avg_cfs_rq
3 0.03% usleep
4 0.04% native_read_msr
6 0.07% ctx_sched_in
6 0.07% common_nsleep
6 0.07% visit_groups_merge
6 0.07% flexible_sched_in
7 0.08% native_write_msr
12 0.13% delay_mwaitx
13 0.14% MM_CompactSchemeFixupObject::fixupArrayObject
18 0.20% MM_CompactScheme::fixupObjectSlot
224 2.43% MM_CompactSchemeFixupObject::fixupObject
769 8.34% MM_CompactScheme::fixupSubArea
836 9.07% MM_CompactSchemeFixupObject::fixupMixedObject
7275 78.89% MM_CompactScheme::getForwardingPtr
9222 --Total--
```

JLM Trace

The java inflated monitors records mean that one or more threads are trying to acquire the object-monitor simultaneously.

GETS is the no of time that a thread tries to acquire a lock. As we can see below the high no of GETS along with TIER2 and TIER3 spinning indicates lock contention in the following class.

```
740 0 0 0 0 0 0 0 0 0 [00007FA4F4069140] MM_PacketList: sublists[]._lock
741 0 0 0 0 0 0 0 0 0 [00007FA4F4069298] MM_PacketList: sublists[]._lock
742 0 0 0 0 0 0 0 0 0 [00007FA4F40693F0] MM_PacketList: sublists[]._lock
743 0 0 0 0 0 0 0 0 0 [00007FA4F4069AA8] MM_PacketList: sublists[]._lock
744 0 0 0 0 0 0 0 0 0 [00007FA4F4069C00] MM_PacketList: sublists[]._lock
745 0 0 0 0 0 0 0 0 0 [00007FA4F406A160] MM_PacketList: sublists[]._lock
746 0 0 0 0 0 0 0 0 0 [00007FA4F406A2B8] MM_PacketList: sublists[]._lock
747
748 Java (Inflated) Monitors
749
750 %MISS GETS NONREC SLOW REC TIER2 TIER3 %UTIL AVER_HTM MON-NAME
751
752 0 121731 121731 0 0 783074 5421 0 1658 [00007FA4F4687958] net/adoptopenjdk/bumblebench/examples/SynchronizedMethodBench@00000000E00
17BF8 (Class)
753
754 LEGEND:
755 %MISS : 100 * SLOW / NONREC
756 GETS : Lock Entries
757 NONREC : Non Recursive Gets
758 SLOW : Non Recursive that Block
759 REC : Recursive Gets
760 TIER2 : SMP: Total try-enter spin loop cnt (middle for 3 tier)
761 TIER3 : SMP: Total yield spin loop cnt (outer for 3 tier)
762 %UTIL : 100 * Hold-Time / Total-Time
763 AVER-HT : Hold-Time / NONREC
```

Either of these tools cannot locate to the method that is being contended by multiple threads.



Call Stack Traces:

- This is used for call graph profiling
- It has a mode for profiling based on contended monitor entry
- Monitors that are going through massive levels of spin cs can record the call stacks

There's a tool called **mxview.jar**

this tool is a java utility for viewing scs output.

- The agent lib command is modified for this particular operation:

agentlib:jprof=scs=monitor_contended_entered+

- The jprof is told to run in a different mode, set itself up for call stack sampling, and monitor contended entered event is told to use.
- The plus sign is there to **record the type of the object** that's being contended.



agentlib: jprof -jar : uses this command for recording jlm and perf data

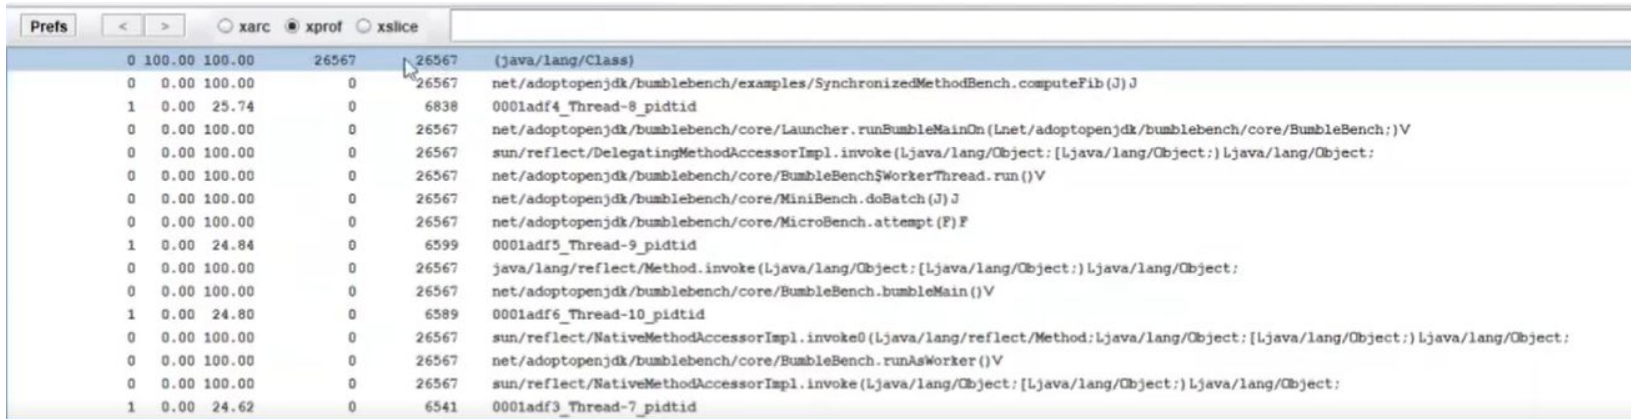
agentlib: jprof=scs=monitor_contended_entered+ : uses this command to record the call stack traces

```
[root@animal training]# /home/ajcraik/sdk/bin/java -DBumbleBench.parallelInstances=4  
-agentlib:jprof=scs=monitor_contended_entered+,logpath=$PWD -jar ./BumbleBench.jar SynchronousMethodBench
```

Rtdriver is used for the operation:

rtdriver -a 127.0.0.1 -c start 15 -c stop 20

After the rtdriver is done recording we can see the **log-rt.1_111633** which has the actual scs output but is not human readable.



lock	time	PID	thread
0	100.00	100.00	26567 (java/lang/Class)
0	0.00	100.00	0 26567 net/adoptopenjdk/bumblebench/examples/SynchronizedMethodBench.computeFib(J)J
1	0.00	25.74	0 6838 0001adf4_Thread-8_pidtid
0	0.00	100.00	0 26567 net/adoptopenjdk/bumblebench/core/Launcher.runBumbleMainOn(Lnet/adoptopenjdk/bumblebench/core/BumbleBench;)V
0	0.00	100.00	0 26567 sun/reflect/DelegatingMethodAccessorImpl.invoke(Ljava/lang/Object:[Ljava/lang/Object;)Ljava/lang/Object;
0	0.00	100.00	0 26567 net/adoptopenjdk/bumblebench/core/BumbleBench\$WorkerThread.run()V
0	0.00	100.00	0 26567 net/adoptopenjdk/bumblebench/core/MiniBench.doBatch(J)J
0	0.00	100.00	0 26567 net/adoptopenjdk/bumblebench/core/MicroBench.attempt(F)F
1	0.00	24.84	0 6599 0001adf5_Thread-9_pidtid
0	0.00	100.00	0 26567 java/lang/reflect/Method.invoke(Ljava/lang/Object:[Ljava/lang/Object;)Ljava/lang/Object;
0	0.00	100.00	0 26567 net/adoptopenjdk/bumblebench/core/BumbleBench.bumbleMain()V
1	0.00	24.80	0 6589 0001adf6_Thread-10_pidtid
0	0.00	100.00	0 26567 sun/reflect/NativeMethodAccessorImpl.invoke0(Ljava/lang/reflect/Method;Ljava/lang/Object:[Ljava/lang/Object;)Ljava/lang/Object;
0	0.00	100.00	0 26567 net/adoptopenjdk/bumblebench/core/BumbleBench.runAsWorker()V
0	0.00	100.00	0 26567 sun/reflect/NativeMethodAccessorImpl.invoke(Ljava/lang/Object:[Ljava/lang/Object;)Ljava/lang/Object;
1	0.00	24.62	0 6541 0001adf3_Thread-7_pidtid

The first value from the list **(java/lang/Class)** is the type of an object that's being locked. There's only 1 lock in the system and there are above 26K locking events.

<div> <div>Prefs</div> <div>< ></div> <div> <input checked="" type="radio"/> xarc <input type="radio"/> xprof <input type="radio"/> xslice </div> </div>						
Parent	0	0.00	100.00	0	26567	net/adoptopenjdk/bumblebench/core/MiniBench.doBatch(J)J
Self	0	0.00	100.00	0	26567	net/adoptopenjdk/bumblebench/examples/SynchronizedMethodBench.computeFib(J)J
Child	0	100.00	100.00	26567	26567	(java/lang/Class)

We can view the **xarc** view of the contended object.

We can see that the object was locked from the synchronized method and can locate the parent of each class.

<div> <div>Prefs</div> <div>< ></div> <div> <input type="radio"/> xarc <input type="radio"/> xprof <input checked="" type="radio"/> xslice </div> </div>						
(java/lang/Class)						
java/lang/reflect/Method.invoke(Ljava/lang/Object;[Ljava/lang/Object;)Ljava/lang/Object;						
net/adoptopenjdk/bumblebench/core/BumbleBench\$WorkerThread.run()V						
net/adoptopenjdk/bumblebench/core/BumbleBench.bumbleMain()V						
net/adoptopenjdk/bumblebench/core/BumbleBench.runAsWorker()V						
net/adoptopenjdk/bumblebench/core/Launcher.runBumbleMainOn(Lnet/adoptopenjdk/bumblebench/core/BumbleBench;)V						
net/adoptopenjdk/bumblebench/core/MicroBench.attempt(F)F						
net/adoptopenjdk/bumblebench/core/MiniBench.doBatch(J)J						
net/adoptopenjdk/bumblebench/examples/SynchronizedMethodBench.computeFib(J)J						
sun/reflect/DelegatingMethodAccessorImpl.invoke(Ljava/lang/Object;[Ljava/lang/Object;)Ljava/lang/Object;						
sun/reflect/NativeMethodAccessorImpl.invoke(Ljava/lang/Object;[Ljava/lang/Object;)Ljava/lang/Object;						
sun/reflect/NativeMethodAccessorImpl.invoke(Ljava/lang/reflect/Method;Ljava/lang/Object;[Ljava/lang/Object;)Ljava/lang/Object;						

We can search from **xslice** for parentheses to find contended objects.

```

package net.adoptopenjdk.bumblebench.examples;

import net.adoptopenjdk.bumblebench.core.Minibench;

public final class SynchronizedMethodBench extends MiniBench {

    public static long next;

    public static long computeFib(long num) {
        long prev = 0, curr = 1;
        while (num-- > 0) {
            synchronized (SynchronizedMethodBench.class) {
                next = prev + curr;
                prev = curr;
                curr = next;
            }
            if (pauseDuration > 0) {
                try {
                    Thread.sleep(0, pauseDuration);
                } catch (InterruptedException e) {
                }
            }
        }
        return prev;
    }

    // Illustrates a MicroBench with an infinite score, for testing purposes

    private volatile long result;

    protected int maxIterationsPerLoop() { return maxIterations; }

    protected long doBatch(long numLoops, int numIterationsPerLoop) throws InterruptedException {
        startTimer();
        for (long i = 0; i < numLoops; ++i) {
            result = computeFib(numIterationsPerLoop);
        }
        pauseTimer();
        return numLoops * numIterationsPerLoop;
    }

    private static int maxIterations;
    private static int pauseDuration;
    static {
        String iters = System.getProperty("SynchronizedMethodBench.iterationCount");
        maxIterations = 9072;
    }
}

```

Contended object found:
java/lang/class



Parent:
net/adoptopenjdk/bumblebench/Synchr
onisedMethodBench.computeFib

- We can figure out the name of the method here: **computeFib**



Parent:
net/adoptopenjdk/bumblebench/core/Minib
ench

- The class SynchronizedMethodBench was extended from Minibench.


```

LV      EVENT  NAME
0        0    0001b464_Thread-7
1        0    net/adoptopenjdk/bumblebench/core/BumbleBench$WorkerThread.run()
V
2        0    net/adoptopenjdk/bumblebench/core/Launcher.runBumbleMainOn(Lnet/
adoptopenjdk/bumblebench/core/BumbleBench;)V
3        0    java/lang/reflect/Method.invoke(Ljava/lang/Object;[Ljava/lang/Ob
ject;)Ljava/lang/Object;
4        0    sun/reflect/DelegatingMethodAccessorImpl.invoke(Ljava/lang/Objec
t;[Ljava/lang/Object;)Ljava/lang/Object;
5        0    sun/reflect/NativeMethodAccessorImpl.invoke(Ljava/lang/Object;[L
java/lang/Object;)Ljava/lang/Object;
6        0    sun/reflect/NativeMethodAccessorImpl.invoke0(Ljava/lang/reflect/
Method;Ljava/lang/Object;[Ljava/lang/Object;)Ljava/lang/Object;
7        0    net/adoptopenjdk/bumblebench/core/BumbleBench.bumbleMain()V
8        0    net/adoptopenjdk/bumblebench/core/BumbleBench.runAsWorker()V
9        0    net/adoptopenjdk/bumblebench/core/MicroBench.attempt(F)F
10       0    net/adoptopenjdk/bumblebench/core/MiniBench.doBatch(J)J
11       0    net/adoptopenjdk/bumblebench/examples/SynchronizedMethodBench.do
Batch(JI)J

```

```

74       0    0001b467_Thread-10
75       1        0    net/adoptopenjdk/bumblebench/core/BumbleBench$WorkerThread.run()
V
76       2        0    net/adoptopenjdk/bumblebench/core/Launcher.runBumbleMainOn(Lnet/
adoptopenjdk/bumblebench/core/BumbleBench;)V
77       3        0    java/lang/reflect/Method.invoke(Ljava/lang/Object;[Ljava/lang/Ob
ject;)Ljava/lang/Object;
78       4        0    sun/reflect/DelegatingMethodAccessorImpl.invoke(Ljava/lang/Objec
t;[Ljava/lang/Object;)Ljava/lang/Object;
79       5        0    sun/reflect/NativeMethodAccessorImpl.invoke(Ljava/lang/Object;[L
java/lang/Object;)Ljava/lang/Object;
80       6        0    sun/reflect/NativeMethodAccessorImpl.invoke0(Ljava/lang/reflect/
Method;Ljava/lang/Object;[Ljava/lang/Object;)Ljava/lang/Object;
81       7        0    net/adoptopenjdk/bumblebench/core/BumbleBench.bumbleMain()V
82       8        0    net/adoptopenjdk/bumblebench/core/BumbleBench.runAsWorker()V
83       9        0    net/adoptopenjdk/bumblebench/core/MicroBench.attempt(F)F
84      10        0    net/adoptopenjdk/bumblebench/core/MiniBench.doBatch(J)J
85      11        0    net/adoptopenjdk/bumblebench/examples/SynchronizedMethodBench.do
Batch(JI)J
86      12        0    net/adoptopenjdk/bumblebench/examples/SynchronizedMethodBench.co
mputeFib(J)J
87      13      7545    (java/lang/Class)

```

Log-rt data for computeFib

```

Java (Inflated) Monitors
%MISS      GETS      NONREC      SLOW      REC      TIER2      TIER3 %UTIL AVER_HTM  MON-NAME
      0    121731    121731         0         0    783074      5421         0      1658  [00007FA4F4687958] net/adoptopenjdk/bumblebench/examples/SynchronizedMethodBench@00000000E00
17BF8 (Class)

LEGEND:
%MISS : 100 * SLOW / NONREC
GETS : Lock Entries
NONREC : Non Recursive Gets
SLOW : Non Recursives that Block
REC : Recursive Gets
TIER2 : SMP: Total try-enter spin loop cnt (middle for 3 tier)
TIER3 : SMP: Total yield spin loop cnt (outer for 3 tier)
%UTIL : 100 * Hold-Time / Total-Time
AVER-HT : Hold-Time / NONREC

```

Jlm data for computeFib example.

```

LV      EVENT  NAME
0       0      000310c7_Thread-4
1       0      Hot.run()V
2       0      Hot.retrieveData(I)Ljava/lang/Object;
3       42     (java/lang/Object)
0       0      000310c8_Thread-5
1       0      Hot.run()V
2       0      Hot.retrieveData(I)Ljava/lang/Object;
3       55     (java/lang/Object)
0       0      000310c9_Thread-6
1       0      Hot.run()V
2       0      Hot.retrieveData(I)Ljava/lang/Object;
3       47     (java/lang/Object)
0       0      000310ca_Thread-7
1       0      Hot.run()V
2       0      Hot.retrieveData(I)Ljava/lang/Object;
3       47     (java/lang/Object)

```

Hot_1

Java (Inflated) Monitors

%MISS	GETS	NONREC	SLOW	REC	TIER2	TIER3	%UTIL	AVER_HTM	MON-NAME
0	184751	184751	28	0	72647767	2185191	100	205012	[00007F9ECC0029B8] java/lang/Object@000000007FFF686I
0	0	0	0	0	0	0	0	0	[00007F9ECC001E08] java/lang/J9VMInternals\$ClassIn:

```

LV    EVENT    NAME
0      0    00039571_Thread-4
1      0    Hot.run()V
2      0    Hot.calculateC()I
3    2305    (java/lang/Object)
3    2304    (java/util/ArrayList)
2      0    Hot.calculateB()I
3      1    (java/util/ArrayList)
0      0    00039572_Thread-5
1      0    Hot.run()V
2      0    Hot.calculateC()I
3    2305    (java/lang/Object)
3    2302    (java/util/ArrayList)
2      0    Hot.calculateB()I
3      1    (java/util/ArrayList)
0      0    00039573_Thread-6
1      0    Hot.run()V
2      0    Hot.calculateC()I
3    2304    (java/lang/Object)
3    2304    (java/util/ArrayList)
2      0    Hot.calculateB()I
3      1    (java/util/ArrayList)
0      0    00039574_Thread-7
1      0    Hot.run()V
2      0    Hot.calculateC()I
3    2305    (java/util/ArrayList)
3    2304    (java/lang/Object)

```

Hot_2

Java (Inflated) Monitors

%MISS	GETS	NONREC	SLOW	REC	TIER2	TIER3	%UTIL	AVER_HTM	MON-NAME
3	79363	79363	2344	0	65023034	1991308	99	475804	[00007F59B4003358] java/lang/Object@000000060A25E1!
2	9842	9842	172	0	4092122	124311	4	166084	[00007F59B4002D28] java/util/ArrayList@000000060A2!

LV	EVENT	NAME
0	0	000310c7_Thread-4
1	0	Hot.run()V
2	0	Hot.retrieveData(I)Ljava/lang/Object;
3	42	(java/lang/Object)
0	0	000310c8_Thread-5
1	0	Hot.run()V
2	0	Hot.retrieveData(I)Ljava/lang/Object;
3	55	(java/lang/Object)
0	0	000310c9_Thread-6
1	0	Hot.run()V
2	0	Hot.retrieveData(I)Ljava/lang/Object;
3	47	(java/lang/Object)
0	0	000310ca_Thread-7
1	0	Hot.run()V
2	0	Hot.retrieveData(I)Ljava/lang/Object;
3	47	(java/lang/Object)

```

public void run(){
    while(true){
        retrieveData(0);
    }
}

public Object retrieveData(int ID) {
    try{
        long threadId = Thread.currentThread().getId();
        synchronized (dataBase) {
            if(sleepType == 0){
                Thread.sleep(sleepTime);
            }else {
                Thread.sleep(0, sleepTime);
            }
        }
        System.out.println("Task done from thread : " + threadId);
    }
}

```

- The contented monitor (java/lang/Object)
- The method name of the synchronized region (retrieveData)
- Stack trace of all the methods called

LV	EVENT	NAME
0	0	00039571_Thread-4
1	0	Hot.run()V
2	0	Hot.calculateC()I
3	2305	(java/lang/Object)
3	2304	(java/util/ArrayList)
2	0	Hot.calculateB()I
3	1	(java/util/ArrayList)
0	0	00039572_Thread-5
1	0	Hot.run()V
2	0	Hot.calculateC()I
3	2305	(java/lang/Object)
3	2302	(java/util/ArrayList)
2	0	Hot.calculateB()I
3	1	(java/util/ArrayList)
0	0	00039573_Thread-6
1	0	Hot.run()V
2	0	Hot.calculateC()I
3	2304	(java/lang/Object)
3	2304	(java/util/ArrayList)
2	0	Hot.calculateB()I
3	1	(java/util/ArrayList)
0	0	00039574_Thread-7
1	0	Hot.run()V
2	0	Hot.calculateC()I
3	2305	(java/util/ArrayList)
3	2304	(java/lang/Object)

```
public int calculateC() {
    try{
        long threadId = Thread.currentThread().getId();
        synchronized (A) {
            synchronized (B) {
                if(sleepType==0){
                    Thread.sleep(sleepTime);
                } else{
                    Thread.sleep(0, sleepTime);
                }
            }
            System.out.println("Task done from thread : " + threadId);
        }
    }
}
```

```
public int calculateB() {
    try{
        long threadId = Thread.currentThread().getId();
        synchronized (B) {
            if(sleepType==0){
                Thread.sleep(sleepTime);
            } else{
                Thread.sleep(0, sleepTime);
            }
        }
        System.out.println("Task done from thread : " + threadId);
    }
}
```

- Contention in 2 methods and 2 types of monitors.
- CalculateC has nested synchronized regions.

LV	EVENT	NAME
0	0	001256d2_Thread-4
1	0	OverlySplit.run()V
2	0	OverlySplit.doSomething()V
3	1726	(java/lang/Object)
3	646	(java/util/ArrayList)
0	0	001256d3_Thread-5
1	0	OverlySplit.run()V
2	0	OverlySplit.doSomething()V
3	658	(java/util/ArrayList)
3	1763	(java/lang/Object)
0	0	001256d4_Thread-6
1	0	OverlySplit.run()V
2	0	OverlySplit.doSomething()V
3	1736	(java/lang/Object)
3	612	(java/util/ArrayList)
0	0	001256d5_Thread-7
1	0	OverlySplit.run()V
2	0	OverlySplit.doSomething()V
3	1771	(java/lang/Object)
3	672	(java/util/ArrayList)
0	0	001256d6_Thread-8
1	0	OverlySplit.run()V
2	0	OverlySplit.doSomething()V
3	1749	(java/lang/Object)
3	637	(java/util/ArrayList)

```

public void doSomething() {
    try{
        synchronized (A) {
            if(sleeptype == 0){
                Thread.sleep(sleepTime);
            }else {
                Thread.sleep(0, sleepTime);
            }
        }
        synchronized (B) {
            if(sleeptype == 0){
                Thread.sleep(sleepTime);
            }else {
                Thread.sleep(0, sleepTime);
            }
            if(sleeptype == 0){
                Thread.sleep(sleepTime);
            }else {
                Thread.sleep(0, sleepTime);
            }
        }
        synchronized (A) {
            if(sleeptype == 0){
                Thread.sleep(sleepTime);
            }else {
                Thread.sleep(0, sleepTime);
            }
        }
    }
}

```



Limitations:

- We cannot locate the java monitor that's being contended by multiple threads.
- If multiple methods are locked, you can see multiple methods here. We can follow up and down the call graph to see where a lock is being acquired and the path that led to it.
- But we don't know if it can locate if the method is called multiple times within the same class.
- We need a parser to read the log-rt file for better understanding.